

Building, Benchmarking, and Deploying an Efficient Spelling Corrector

NLP Group Assignment — Question 3

Group 21: Krishna Jhanwar (12341250), Siddharth Rai (12342050), Slok
Tulsyan (12342080), Vasu Garg (12342330), Gaurav Kumar (12340790)

September 8, 2026

Abstract

This report documents the design, implementation, and evaluation of a spelling corrector built on the NLTK Brown corpus, capable of correcting both *non-word* errors (out-of-vocabulary misspellings) and *real-word* errors (contextually incorrect but valid words) within an edit distance of one. Two independent candidate-generation strategies are implemented and compared: a brute-force edit-distance-1 generator, and the Symmetric Delete Spelling Correction (SymSpell) algorithm. The two are benchmarked head-to-head on a fixed batch of 1,000 misspelled words, showing a $\sim 5.8\times$ speed-up for the Symmetric Delete method. The corrector is evaluated for accuracy on an automatically generated, edit-distance-1 test set drawn from 10% of the Brown corpus sentences, and finally deployed as a live, interactive terminal application.

1 Introduction

Spelling correction is one of the oldest and most practically important problems in natural language processing. A useful corrector must handle two qualitatively different kinds of mistakes:

- **Non-word errors** — the misspelled string is not a valid word at all (e.g. “*sentnce*”), so any out-of-vocabulary token is immediately flagged and the job is to rank the plausible corrections.
- **Real-word errors** — the misspelled string happens to collide with another valid word (e.g. “*sea*” instead of “*see*”), so the word cannot be flagged from vocabulary membership alone; only the surrounding context reveals the mistake.

This assignment asks us to build both capabilities on top of the NLTK **Brown corpus**, compare two ways of generating single-edit candidate corrections, benchmark them for speed, and package the result as a live CLI tool. Section 2 covers corpus and model preparation, Sections 3 and 4 cover candidate generation and correction logic, Section 5 covers evaluation and the “Speed Demon” benchmark, Section 6 covers the interactive application, and Section 7 discusses results, limitations, and possible extensions. Full source code is reproduced in the Appendix.

1.1 System overview

Figure 1 sketches the overall data flow: the Brown corpus is processed once into a vocabulary, a unigram frequency table, and a bigram language model; these are shared by both candidate-generation methods and by the correction logic that sits on top of them; the same corrector object then powers both the offline evaluation harness and the live CLI.




Figure 1: End-to-end pipeline of the spelling corrector.

2 Part 1: Corpus and Model Preparation

2.1 Vocabulary and unigram frequencies

The Brown corpus is loaded via `nltk.corpus.brown.sents()`. Each sentence is lower-cased and filtered to purely alphabetic tokens, since numbers and punctuation have no meaningful “edit distance 1” correction. Word counts are accumulated with a `collections.Counter`:

$$\text{freq}(w) = \#\{\text{occurrences of } w \text{ in the corpus}\}, \quad V = \{w : \text{freq}(w) > 0\}$$

2.2 Bigram language model

A bigram model estimates $P(w_2 | w_1)$ from corpus counts, with sentence boundaries marked by pseudo-tokens `<s>` and `</s>` so that sentence-initial and sentence-final words are also modeled. Raw maximum likelihood estimates,

$$P_{\text{MLE}}(w_2 | w_1) = \frac{C(w_1, w_2)}{C(w_1)},$$

assign zero probability to any bigram never seen in training, which is fatal when comparing candidate phrases that are individually rare. We therefore use **add-one (Laplace) smoothing**:

$$P(w_2 | w_1) = \frac{C(w_1, w_2) + 1}{C(w_1) + |V|}$$

where $|V|$ is the vocabulary size (plus the two pseudo-tokens). Phrase probabilities are computed in **log-space** to avoid numerical underflow and to make comparisons between candidate phrases a simple subtraction:

$$\log P(w_1 \dots w_n) = \sum_{i=1}^n \log P(w_i | w_{i-1})$$

2.3 Caching

Because building these tables from 56K+ sentences is the most expensive one-time step, the vocabulary, frequency table, bigram counts, and sentence list are pickled to `data/*.pkl` after the first run, so all later runs (evaluation, demo, CLI) load instantly instead of rebuilding.

3 Part 2: Candidate Generation Methods

Both methods below are meant to answer the same question — “*which vocabulary words are within one edit of this string?*” — but they arrive at the answer through very different computational routes, which is exactly what the Speed Demon benchmark in Section 5 is designed to expose.

3.1 Method A: Standard Edit-Distance-1 Generation

Method A is the classical brute-force generator (in the style popularised by Norvig’s spelling corrector). For a word of length n it enumerates every string reachable by a single *delete*, *transpose*, *replace*, or *insert* operation:

$$\begin{aligned} \text{deletes} &: L + R[1:] \text{ for each split } (L, R) \\ \text{transposes} &: L + R[1] + R[0] + R[2:] \text{ for each split with } |R| > 1 \\ \text{replaces} &: L + c + R[1:] \text{ for each split, each letter } c \\ \text{inserts} &: L + c + R \text{ for each split, each letter } c \end{aligned}$$

This produces on the order of $26n + 25$ candidate strings, most of which are not real words; the raw set is intersected with the vocabulary to keep only valid candidates.

Complexity: $O(26n)$ string constructions per query word, each followed by an $O(1)$ average-case hash-set membership check — i.e. $O(n)$ work scaled by the alphabet size.

3.2 Method B: Symmetric Delete Spelling Correction

Symmetric Delete (SymSpell) trades query-time work for a one-off pre-processing cost. **Pre-processing** (done once, offline): for every vocabulary word w , generate all of its single-character deletions and map each deletion string back to w (plus an identity entry $w \mapsto w$):

$$\text{delete_dict}[d] += w \quad \text{for every } d \in \text{deletes}(w) \cup \{w\}$$

Query time: for a misspelled word q , generate *only* its own single-character deletions (plus itself) and look each one up directly in `delete_dict` — no alphabet loop at all:

$$\text{candidates}(q) = \bigcup_{v \in \text{deletes}(q) \cup \{q\}} \text{delete_dict}[v]$$

Because deletion-based matching can, in rare cases, surface a pair of words whose *true* edit distance is 2 (a known property of delete-delete matching), a cheap Levenshtein-distance ≤ 1 verification pass is applied to the (small) candidate set before it is returned, so Method B stays exactly as correct as Method A while remaining fast.

Complexity: pre-processing is $O(n \cdot |V|)$, paid once; each query is $O(n)$ deletions plus $O(n)$ hash lookups plus $O(k)$ verification where k is the (small) number of raw hits — i.e. *no* factor of 26 on the query-time hot path.

3.3 Summary comparison

	Method A (Brute force)	Method B (Symmetric Delete)
Pre-processing	None	$O(n \cdot V)$, once
Per-query cost	$O(26n)$	$O(n)$
Operations at query time	delete, transpose, replace, insert	delete only
Extra vocabulary filter needed	Yes (membership check)	Verification only

Table 1: Method A vs. Method B, computational comparison.

4 Part 3: Spelling Correction Logic

4.1 Non-word error correction

For a token $q \notin V$, candidates are pooled from *both* Method A and Method B (their union), and the corrector picks the candidate with the highest unigram frequency:

$$\hat{w} = \arg \max_{c \in \text{cand}_A(q) \cup \text{cand}_B(q)} \text{freq}(c)$$

If no candidate is found within edit distance 1, the original (unrecognized) word is returned unchanged.

4.2 Real-word error correction

For a token $w \in V$ that may nonetheless be wrong in context, the corrector still generates edit-distance-1 candidates for w (the same V -filtered union of Method A/B), but instead of ranking by raw frequency it compares the bigram log-probability of the local phrase — the target word together with its immediate left and right neighbours — with w in place versus each candidate c in place:

$$\Delta(c) = \log P(\text{phrase with } c) - \log P(\text{phrase with } w)$$

A candidate is only accepted if it beats the original by a fixed margin,

$$\hat{w} = \begin{cases} c^* & \text{if } \Delta(c^*) > \tau, \ c^* = \arg \max_c \Delta(c) \\ w & \text{otherwise} \end{cases}$$

with $\tau = 1.5$ nats in our implementation (`REAL_WORD_LOG_MARGIN`). This margin exists because bigram counts over a moderately sized corpus like Brown are sparse; without a margin, small, noisy probability differences would cause the corrector to “fix” words that were never wrong.

4.3 Full-sentence correction

`correct_sentence` ties both cases together: every token in an input sentence is routed to `correct_nonword` (if out-of-vocabulary) or `correct_realword` (if in-vocabulary), and the sentence is reassembled with original spacing/punctuation preserved around unchanged tokens. This is the function the live CLI (Section 6) calls directly.

5 Part 4: Evaluation and “Speed Demon” Benchmark

5.1 Test set generation

10% of Brown corpus sentences are randomly sampled (fixed seed for reproducibility). For each sampled sentence, one word of length ≥ 3 is chosen at random, and a single random edit operation (delete/insert/replace/transpose, uniformly chosen) is applied, retrying until the result:

- is *not* in the vocabulary $\Rightarrow$ a non-word-error example, or
- is a *different* vocabulary word $\Rightarrow$ a real-word-error example.

This produced **5,658** non-word examples and **3,408** real-word examples (not every sentence yields both, since not every word has a valid single-edit neighbour in the vocabulary).

5.2 Accuracy

Accuracy is the fraction of test examples for which the corrector’s output exactly matches the original (pre-corruption) word.

Test set	Correct	Total	Accuracy
Non-word errors	4,745	5,658	83.8%
Real-word errors	2,293	3,408	67.3%

Table 2: Correction accuracy on the generated test set.

Non-word correction is markedly more accurate than real-word correction, which is expected: a non-word error gives an unambiguous, immediate signal (vocabulary membership) and unigram frequency alone is a strong prior over plausible English words. Real-word correction, in contrast, must decide *whether to act at all* based on a much sparser and noisier signal (bigram counts over a 3-word window), and the conservative margin τ deliberately trades some recall for precision — it declines to “fix” a real word unless the context strongly favours a swap.

5.3 Speed Demon benchmark

The non-word candidate-generation step is isolated from the rest of the pipeline (frequency ranking, bigram scoring) and timed in a controlled, head-to-head setup: the *exact same* batch of 1,000 misspelled words is passed first through Method A’s generator, then through Method B’s, using `time.perf_counter()` for wall-clock timing.




Figure 2: Non-word vs. real-word correction accuracy.

Method	Total time (1,000 words)	Speed-up
A — Brute-force edit distance 1	0.1126 s	1× (baseline)
B — Symmetric Delete	0.0193 s	5.8×

Table 3: Speed Demon benchmark results.

Why Method B wins. Method A reconstructs the full set of edit-distance-1 strings *from scratch* for every query, an $O(26n)$ operation dominated by the alphabet loop inside the *replace* and *insert* branches. Method B never performs an alphabet loop at query time at all: it only ever deletes characters from the query word ($O(n)$ deletions) and performs $O(n)$ hash lookups against a dictionary that was built once, offline, by pre-computing deletions of every vocabulary word. In other words, Method A pays the full combinatorial cost of relating a word to its neighbourhood on *every query*; Method B pays a similar combinatorial cost only *once*, during pre-processing, and reduces every subsequent query to cheap deletions and hash lookups. This is precisely the $\sim 5.8\times$ speed-up observed empirically, and the gap would widen further for longer words or larger query batches, since Method A’s per-query cost scales with n while Method B’s advantage compounds with reuse of the same pre-processed dictionary across queries.

6 Part 5: Live Interactive Application

The corrector is packaged as a continuous terminal CLI (`src/cli.py`): a `while True` loop reads one sentence at a time, corrects it via `correct_sentence`, prints the corrected sentence with any changed words wrapped in `**asterisks**` (and, on ANSI-capable terminals, also rendered in bold green), and prints the wall-clock latency of the correction in milliseconds. The loop terminates when the user types `exit`.




Figure 3: Total latency for 1,000 misspelled words, Method A vs. Method B.

Listing 1: Sample interactive session

```

1 =====
2 Spelling Corrector -- interactive mode
3 Type a sentence and press Enter. Type 'exit' to quit.
4 =====
5
6 > I hav a good feeling about this.
7 Corrected: I **had** a good feeling about this.
8 Changed:   hav -> had
9 Latency:   1.45 ms
10
11 > This is a test sentnce.
12 Corrected: This is a test **sentence**.
13 Changed:   sentnce -> sentence
14 Latency:   0.83 ms
15
16 > I would like to sea the world.
17 Corrected: I would like to **see** the world.
18 Changed:   sea -> see
19 Latency:   2.10 ms
20
21 > exit
22 Goodbye!

```

Note the third example: “sea” is a perfectly valid English word, so it is never flagged by vocabulary lookup alone; it is only corrected because the real-word pathway (Section 4) finds that “I would like to see the world” is significantly more probable under the bigram model than the original phrase.

7 Discussion, Limitations, and Extensions

Non-word vs. real-word accuracy gap. The ~17-point accuracy gap between non-word (83.8%) and real-word (67.3%) correction is consistent with the wider spelling-correction literature: context-dependent decisions are inherently harder than vocabulary-membership decisions,

and a bigram model is a comparatively weak language model (it only “sees” one word of context on each side).

Smoothing choice. Add-one (Laplace) smoothing was chosen for simplicity and interpretability; it is known to over-penalize frequent bigrams in favour of unseen ones on larger vocabularies. A production system would likely switch to a more sophisticated scheme (e.g. Kneser–Ney or Katz backoff) or a higher-order (trigram+) model, at the cost of additional implementation complexity.

The real-word margin τ . This single hyperparameter directly trades recall for precision on real-word correction: lowering τ would catch more genuine errors like “meat” $\rightarrow$ “meet” but would also risk more false-positive “corrections” of already-correct sentences. Tuning τ against a labeled validation set (rather than the fixed value used here) is a natural next step.

Scope. Both candidate-generation methods are explicitly scoped to edit distance 1, per the assignment specification. Multi-error words (edit distance ≥ 2) are out of scope and would require extending Method B’s pre-processing to deeper deletion depths (which SymSpell natively supports) and Method A to a recursive edit-distance-2 generator.

Reproducibility. All random sampling (test-set generation, Speed Demon batch construction) is seeded, so the numbers reported in Section 5 are exactly reproducible by re-running `python main.py evaluate`.

8 Conclusion

This assignment implements a complete, working spelling correction pipeline on top of the NLTK Brown corpus: a unigram/bigram language model (Part 1), two independent single-edit candidate generators with markedly different computational profiles (Part 2), frequency- and context-driven correction logic for non-word and real-word errors respectively (Part 3), an automated evaluation harness with an explicit head-to-head speed benchmark (Part 4), and a live, interactive terminal deployment (Part 5). The central empirical result — Method B (Symmetric Delete) outperforming Method A (brute-force edit distance 1) by $\sim 5.8\times$ on an identical workload — is a direct, measurable consequence of moving the alphabet-sized combinatorial search from query time into a one-time pre-processing step, illustrating a classic and broadly applicable systems-design trade-off between pre-computation and per-query work.